



Machine Learning BATCH 8

Deepthishree G S 18Z312

Harshini S 18Z320

Iswaryaa GP 18Z323


Janani R 18Z324

Swetha M 18Z360

Topic: Reinforcement
Learning



Problem Statement



- Topic: **Dots and Boxes**
- To train an agent to play a simple pen and pencil game **Dots and Boxes** using a reinforcement learning (Q-learning Algorithm)
- To make the agent learn the game from the rewards and penalties
- To check its performance by playing against human and random agents
- To deepen the understanding of Q Learning algorithm by building it from scratch



Why RL for this problem?

- There is no training data available beforehand
- The agent improves by playing more and more games
- The more the games, the better the agent learns from the rewards and penalties from environment
- Since, the agent learns the best move to from boxes based on the rewards, reinforcement learning will be best for the problem



Problem Design

Task: Play Dots and Boxes

Performance: Win-rate against itself, humans and a random agent

Training Experience: Play 10000 Games against itself

Target Function: Policy : State \rightarrow Action

Target Function representation:

Q Table $Q(s,a) = \text{Immediate Reward} + \text{DR} * V(\text{Next State})$

Policy is to choose the action with maximum Q value for the given state

(Reference for Equations: Tom Mitchell Ch 13)



Assumptions and representations

- 3x3 grid with 4 boxes. 2 player game
- Player with greater number of boxes wins. There is a possibility of tie (unlike a few variations of the game)
- Assign Line numbers from 1 to 12 to accept input and represent the board state.
- Board state is represented as bit array.

Eg: 000000000000 is the initial state and 010101010101 is an intermediary state with 6 lines drawn

- Action is a number from 1 to 12 indicating legal moves from a state.

Eg: in the above mentioned state 1, 12 possible actions {1, 2, 3, ... 12} and state 2, only six possible actions {1, 3, 5, 7, 9, 11}



Board States

| | | | | |
|----|---|----|---|----|
| . | 1 | . | 2 | . |
| 7 | | 8 | | 9 |
| . | 3 | . | 4 | . |
| 10 | | 11 | | 12 |
| . | 5 | . | 6 | . |

State: 000000000000

Initial board state

| | | | | |
|----|----|----|----|----|
| . | 1 | . | 2 | . |
| 7 | | | | 9 |
| . | -- | . | -- | . |
| 10 | | 11 | | 12 |
| . | 5 | . | 6 | . |

State: 001100010000

Intermediate board state

| | | | | |
|---|----|---|----|---|
| . | -- | . | -- | . |
| | | | | |
| . | -- | . | -- | . |
| | | | | |
| . | -- | . | -- | . |

State: 111111111111

Final board state



Design Choices

- **Learning Rate**=0.2 [Each update of QTable affects learning slowly]
- **Discount Rate**=0.8 [Importance of future rewards compared to immediate rewards]
- **Reward**: 200 for winner, 100 for each box filled (only 4 boxes need to be filled), 50 for tie (more common in this game)
- **Penalty**: 200 for loser, 100 for each box opponent fills
- **Score**: 2 for each box completed
- **Exploration** (50 %) using random moves and visiting states with least visit count and then **Exploitation** using Q Table



Results and Inference

→ Against Random agent:
(1000 games)

- ◆ Win-rate = 0.921
- ◆ Tie-rate = 0.064
- ◆ Lose-rate = 0.015

→ Against Human:
(10 games)

- ◆ Win-rate = 0.3
- ◆ Tie-rate = 0.5
- ◆ Lose-rate = 0.2

THANK YOU!

